

Curating Research Assets in Behavioral Sciences: A Tutorial on the Git Version Control System

Matti Vuorre¹ & James P. Curley^{1,2}

¹ Department of Psychology, Columbia University, New York, USA

² Department of Psychology, University of Texas at Austin, Texas, USA

Abstract

Recent calls for improving the reproducibility of behavioral sciences have increased attention to the ways in which researchers curate, share and collaborate on their research assets. However, these discussions have largely failed to provide practical instructions on how to do so. In this tutorial paper, we explain how version control systems, such as the popular Git program, address these challenges to reproducibility. We then present a tutorial on how to use Git from the computer's command line, and with the popular graphical interface in the R Studio development environment. This tutorial is especially written for behavioral scientists with no previous experience with version control systems or command line interfaces, and covers single-user and collaborative workflows. Git is easy to learn, presents an elegant solution to specific challenges to reproducibility, facilitates multi-site collaboration and productivity by allowing multiple collaborators to work on the same source simultaneously, and can be implemented on common behavioral science workflows with little extra effort. Git may also offer a suitable solution to transparent data and material sharing through popular online services, such as GitHub and Open Science Framework.

Keywords: reproducibility; version control; git; research methods; open science

Word count: 7700

Introduction

Lack of reproducibility is increasingly recognized as a problem across scientific disciplines, and calls for changing the scientific workflow to enhance reproducibility have been published in a wide range of research areas, including Biology (Markowetz, 2015), Ecology (Ihle, Winney, Krystalli, & Croucher, 2017), Neuroscience (Eglen et al., 2017) and Psychology (Munafò et al., 2017). Some studies suggest that one specific challenge to reproducibility is the ways in which researchers organize, curate, share, and collaborate on the research assets—e.g. data and code used to analyze the data—which support the product’s conclusions (Vanpaemel, Vermorgen, Deriemaeker, & Storms, 2015; Wicherts, Borsboom, Kats, & Molenaar, 2006).

Fortunately for the empirical sciences, challenges related to organizing and curating materials across time, space, and collaborators have been solved to a high standard in computer science with Version Control Systems (VCS). In this tutorial article, we introduce a popular VCS called Git, and illustrate its use in the scientific workflow with a hypothetical example project. We show how to use Git by using text commands from the computer’s command line, and with a Graphical User Interface (GUI) implemented in the popular R Studio Integrated Development Environment (IDE; RStudio Team (2016)). We also show how to use Git with the popular online service GitHub for collaborative workflows. Learning how to use Git(Hub) will streamline workflows and help researchers to stay better organized, and thereby facilitate reproducibility.

Version Control Systems

Before proceeding with the practical tutorial, it is important to understand what Version Control Systems (VCS) are and what problems they were designed to address. VCS are computer programs designed for tracking changes to computer code, and collaborating on the code with others. They were initially developed by software engineers for writing code collaboratively¹, but are increasingly being adopted to enhance workflows outside computer science. To help understand why, it is helpful to think of “code” more broadly as any text written on a computer: Manuscripts, books, statistical analysis scripts, source code of computerized experiments, and even data files are “code” or have source code, which is just plain text written on a computer, and which can benefit from version control. Going forward, when we use the word “code” in this tutorial, we mean it in this broad sense (e.g. this manuscript’s source “code” was written on a computer, and was version controlled.)

For example, computerized behavioral experiments are usually written with computer programming languages such as MATLAB or Python. The experiment’s source code (text written by humans but interpreted by computers to e.g. display stimuli to participants) may have multiple authors and go through multiple versions. Keeping track of the versions of the program and changes to the code, and allowing many authors to contribute to it (without breaking the experimental program) are problems that VCSs were specifically designed for. It might be less obvious that writing a manuscript is quite similar, at least as far as the computer is concerned: Multiple authors write multiple versions of a text document, and sometimes previous versions need to be inspected, and the text needs to be “merged”

¹The software we present below is used by major software developers such as Microsoft, Google, and Facebook on code bases with hundreds of contributors.

49 across the many authors. Even less obvious is that datasets are often plain text: In most
50 computerized behavioral experiments, the output data are numbers and text written into a
51 text file. These text files can be version controlled (and researchers would usually not want
52 to see that their raw data files have changed after they were created—VCS allows verifying
53 that they haven’t changed because their history is logged.)

54 The core concepts of version control are that contributors to a project create small
55 checkpoints of the changes they make to the source code—analogous to saving an intermediate
56 version of a file on the computer’s hard disk—and then submit those changes to the VCS. The
57 VCS maintains a history of changes to the code between these little checkpoints, and therefore
58 allows coming back to any earlier version by browsing the history. These concepts—i.e. the
59 typical Git workflow—is illustrated in Figure 3.

60 To understand what VCS is, it is useful to contrast an example scenario with VCS to
61 what we believe is a fairly standard workflow. Consider collaborating with one person on a
62 manuscript that reports results from a data set using some statistical model. In the standard
63 workflow, one person might format the raw data in one way to fit a statistical model, and
64 then write a draft of the manuscript. This would create three files: `data.csv`², `analysis.R`,
65 and `manuscript.doc`. If the coauthor then wished to explore another statistical model,
66 which requires the data in a different format, and then edit the manuscript, she would create
67 three more files: `data_new.csv`, `analysis_new.R`, and `manuscript_new.doc`. This cycle
68 would then repeat as many times as required, each time creating more files, making it more
69 and more difficult for the authors to remember which data was paired with which analysis,
70 and which analyses (and data format) was reported in which version of the paper.

71 VCS greatly streamlines this workflow: With VCS, there would be only three files
72 (data, analysis, and manuscript), but they would all be under version control, which would
73 keep track of changes to the files and their different versions. Because VCS monitors changes
74 to the files and allows saving a history for each of them, creating new files for every new
75 idea or edit is unnecessary. This lack of duplicity, in turn, possibly reduces room for error
76 in remembering which data file was linked with which analysis, and which manuscript
77 version had the correct numerical results, and so forth. Further, as we explain below, VCS
78 allows many coauthors to work on the files simultaneously, eschewing the need for emailing
79 different versions of the same file to the coauthors, then waiting for them to make changes
80 before continuing with your work. Box 1 shows practical examples of scenarios where VCS
81 (specifically, Git, see below) can be used to improve the scientific workflow. To get a birds-eye
82 view of the Git workflow, before beginning with the details, we suggest readers to inspect
83 Figures 3 and 7 at this point.

².`csv` stands for comma separated values, and is a plain text file commonly used to store two-dimensional data.

Box 1. How Git facilitates the scientific workflow.

How to try different ways of visualizing and modeling data while keeping track of the different versions—and the differences between them?

- Git saves each version of the analysis, allowing testing new features without losing previous versions or proliferating files in the project's directories. All the past versions can be directly compared with each other, and retrieved from the Git's history.

How to work on the same code or manuscript files simultaneously with collaborators?

- Git was designed for distributed collaborative work: Collaborators work on their local copies and "push" and "pull" material to and from each other or a central "repository". Git keeps precise track of who has done what, when, and (possibly) why. Git never loses information or overwrites work with another collaborator's work, but allows for many authors to work on the same ideas simultaneously.

How to share work easily, safely, and in an organized manner?

- Git enforces a common organization scheme among collaborators, making it easier to keep everyone "on the same page" with what goes where, and how to contribute to specific parts of the project. Git projects are shared as a whole, so complex projects with multiple files linking to each other are easy to share. Sharing projects with others is built into Git, and can be facilitated with online services such as GitHub: Once a project uses Git, it can be very easily copied to GitHub, from where others can easily download the entire project onto their local computers.

84

85 **The Git Version Control System**

86 Version control software has a long history in software engineering, and there are many
 87 VCS programs. Some popular ones are Apache Subversion³, Mercurial⁴, and Git⁵. In this
 88 tutorial, we focus on Git because it is already increasing in popularity within the scientific
 89 community, and is especially good for scientific collaboration because of online tools which
 90 allow seamless collaboration even for very large research teams. One of these online tools is
 91 GitHub, which we discuss in more detail below. Further, Git is free and open source, works
 92 on Windows, Mac, and Linux operating systems (among others).

93 Git, unlike the operating system's (OS) default file viewer (Mac's Finder, Windows'
 94 File Explorer), is not a point-and-click program for navigating files and folders on a computer.
 95 Instead, Git adds functionality to the existing file system, by making available a specific set of
 96 commands—either executed from the command line, or through a graphical point-and-click
 97 user interface (GUI)—which allow keeping track of files and their history, and distributing
 98 the files across multiple computers and users. Because Git is a standalone program not
 99 usually included in standard OS installs, users must first download and install the Git
 100 software on their computers.

³<https://subversion.apache.org/>

⁴<https://www.mercurial-scm.org/>

⁵<https://git-scm.com/>. The creator of Git, Linus Torvalds (who also is the principal developer of the Linux operating system), named Git after himself as "the stupid content tracker" (McMillan (2005); <https://git-scm.com/docs/git.html>)

Installing Git

Even if you already have Git installed (some computers do) it is good practice to install the latest version⁶, which can be downloaded as a standalone program from <https://git-scm.com/download>. For Mac users, the easiest way to install or update Git to the latest version is to download the installer from <http://git-scm.com/download/mac>, and install it like any other program. Similarly, Windows users can download the Git software installer from <http://git-scm.com/download/win> and install it like any other application. Linux (and other) users can download Git and find further install instructions on the Git website (<https://git-scm.com/book/en/v2/Getting-Started-Installing-Git>).

Git setup

Once you have installed Git, you can use it from the computer’s command line or through a GUI—there are various GUIs for using Git, and below we will explain how to use the R Studio Git GUI. Before using a GUI, it is important to learn a few basic Git commands from the OSs command line, because it is simultaneously the most basic and most flexible interface to Git’s functionality. Understanding the basic Git commands as entered through the computer’s command line also facilitates understanding its more advanced uses, and is highly recommended.

The command line. The command line is a text-based interface for interacting with your computer, and its functionality greatly extends that of the standard way of interacting with the computer by clicking and pointing with a mouse. Many advanced techniques require using your computer through a command line, and it is very helpful in e.g. scripting and scheduling tasks on your computer. Here, we introduce the command line in just enough detail so that you can navigate folders on the computer, and set up Git’s basic configuration (identify Git’s user).

To access the command line interface, you need to use a command line “shell” application. Mac users can open the built-in app Terminal, and Windows users can use the Git Bash application, which is installed with the Windows Git program. After opening the command line shell, you can type in commands and execute them by pressing Return (Mac) or Enter (Windows). The most common command line functions are listed in Table 1.

First, you need to know how to navigate the folders on your computer (a task that is typically done by clicking folders in Finder (Mac) or File Explorer (Windows)). This is important because whenever you execute commands in the command line, they are executed in a specific directory. Usually, when you open up your command line shell, you begin in the user’s home directory (or folder, we use these terms interchangeably). Depending on your operating system, the home directory is usually represented with a `~` on the left side of the cursor. To ask for the current working directory, you can use the function `pwd`. To move up in the directory hierarchy (into the folder that contains the current directory), you can use `cd ..` (note the space). To move into a folder that is inside the current working directory, you can use `cd folder` where `folder` is the name of the desired folder. These command line functions are illustrated in Figure 1 with the Mac Terminal application.

Some users may find using a text-based command line interface unfamiliar, but to get started with Git, there are two required configuration commands which you need to run

⁶As of the writing of this article, the current version of Git is 2.13.1.

Table 1

Basic command line functions.

Command	Result
<code>pwd</code>	Show current directory.
<code>cd ..</code>	Move out of current directory (up one level).
<code>cd folder</code>	Move into folder (must be inside current directory).
<code>ls</code>	List files and folders in current directory.

Note. Execute the commands in your command line shell application (e.g. Terminal (Mac) or Git Bash (Windows)) by pressing Return (Mac) or Enter (Windows).

```

Matti [~]$ pwd
/Users/Matti
Matti [~]$ cd ..
Matti [/Users]$ pwd
/Users
Matti [/Users]$ cd Matti
Matti [~]$ pwd
/Users/Matti
Matti [~]$

```

Figure 1. The Terminal command line shell. The functions are explained in more detail in the text. Each command (prepending with a \$ symbol) is followed by its output on the following line. To the left of the \$ symbol, this user's shell application displays the user's user name and the current directory in square brackets. Here, the outputs are directories separated with forward slashes, and top-most (containing, also known as parent) folders are on the left of their subfolders (also known as children). Your command line interface might look slightly different because of different user and folder names, operating systems, and command line shell applications.

once, and the basic functionality requires using only four commands⁷. However, users can also read through the first part of the tutorial without executing the commands and simply focus on the workflow; we show how to use Git with a GUI in the second part of the tutorial.

Setting Git's user information. The first step in using Git is making it aware of who is using the computer. You need to set the user's name and email address by entering a few basic commands in the command line. First, to show the current user information, run the following command in the command line:

```
$ git config --global user.name
```

This should not return anything, unless a previous user of the computer has set the global Git user name. Each Git command starts with the word `git`, then a command (such as `config`), and then arguments to the command, such as `--global` (for global configuration),

⁷If using a text-based command line seems challenging, Codecademy (<https://www.codecademy.com/learn/learn-the-command-line>) has a free interactive online tutorial, and MIT offers a free online game to teach using the command line (<http://web.mit.edu/mprat/Public/web/Terminus/Web/main.html>).

```
Matti ~$ git config --global user.name
Matti ~$ git config --global user.name "Matti Vuorre"
Matti ~$ git config --global user.name
Matti Vuorre
Matti ~$ git config --global user.email "mv2521@columbia.edu"
Matti ~$ git config --global user.email
mv2521@columbia.edu
Matti ~$
```

Figure 2. Setting up Git's configuration using the command line. Notice that for these configuration commands, the current working directory does not matter (we did them in the user's home directory).

153 followed by variables, such as `user.name`. To ensure that Git knows who you are, execute
 154 the following command (replacing `User Name` with your desired user name):

```
$ git config --global user.name "User Name"
```

155 This command maps `User Name` to Git's global `user.name` variable. If you now re-run
 156 the first command (`git config --global user.name`), the command line will return the
 157 user name you entered. The user name can be anything you'd like, but it is probably a
 158 good idea to use your real name so that potential collaborators know who you are. The
 159 second piece of information is your email address, which is entered by the following command
 160 (where `email@address.com` is your email address):

```
$ git config --global user.email "email@address.com"
```

161 You can verify that the correct email address was saved with `git config --global`
 162 `user.email`. These commands are shown as entered to the Terminal command line shell
 163 application in Figure 1.

164 Once this information is entered, Git will know who you are, and is able to track who is
 165 doing what and when within a project, which is especially helpful when you are collaborating
 166 with other people, or when you are working on multiple computers. For detailed instructions
 167 on how to use Git commands, you can type `git --help`. For help on how to use the `git`
 168 `config` command, type `git config --help`. When printed in the command line, some
 169 help pages run for several pages; you can press the space bar to move to the next page, or `q`
 170 to quit looking at the help page.

Using Git

171
 172 The first operating principle of Git is that your work is organized into independent
 173 projects, which Git calls *repositories*⁸. A repository is a folder on your computer which is
 174 version controlled by Git⁹. Everything that happens inside a repository is tracked by Git,
 175 but you have full control of what is committed to Git's history and when. Because you have

⁸For advanced users, Git *submodules* allow linking projects to each other, or organizing more complex projects into projects and their sub-projects (<https://git-scm.com/book/en/v2/Git-Tools-Submodules>).

⁹There are no visible changes to a folder once it is tracked by Git. Once Git is initialized in a folder, the only change is that a hidden folder, called `.git` is added, but users do not need to interact with it directly.




Figure 3. A diagram illustrating the typical Git workflow. Verbs indicated in bold text are Git operations and explained in detail in the main text. In brief, this workflow consists of making changes to files in the project (such as editing a manuscript.) Once the revised version of the file is saved, the user adds the changed file to the staging area (1). Many files can be added to the staging area, if desired. Once the changes in the staging area reflect a conceptual entity, such as edits to an image file and its accompanying caption in the manuscript, the user then commits the changes to Git’s history (2). These commits are accompanied with short commit messages that describe the changes made in that commit. Finally, when needed, changes to files can be discarded by checking out an earlier version of the file from Git’s history. Figure adapted from <https://git-scm.com/book/en/v2/Getting-Started-Git-Basics>.

full control of what and when is committed to history, there is a small set of operations you need to know¹⁰.

Briefly, when you work in a Git repository (make changes to files within it), Git monitors the state of each file, and when they change Git knows that they differ from the previously logged state. If you are happy with the current changes, you **add** the changed files to Git’s “staging area”. If you then are certain that the changes in the staging area are desirable, you **commit** the changes. These two operations are the backbone of using Git to store the state of the project whenever meaningful changes are made. Importantly, each commit in the repository’s history contains information to recover the full state of the project at that point in time. Users can always go back to an earlier version by **checking out** a previous state from Git’s history. This core of the Git workflow is illustrate in Figure 3.

To understand the Git workflow in practice, we now turn to a practical example using a hypothetical project. Git can be added to a project at any stage of the project’s life cycle, but to most clearly show its use, we begin with an empty project with nothing in it.

¹⁰It helps to have a Git command cheat sheet (<https://services.github.com/resources/cheatsheets/>) printed and taped on your wall, but it contains many more commands than are needed for the basic use of Git in standard Psychology studies.

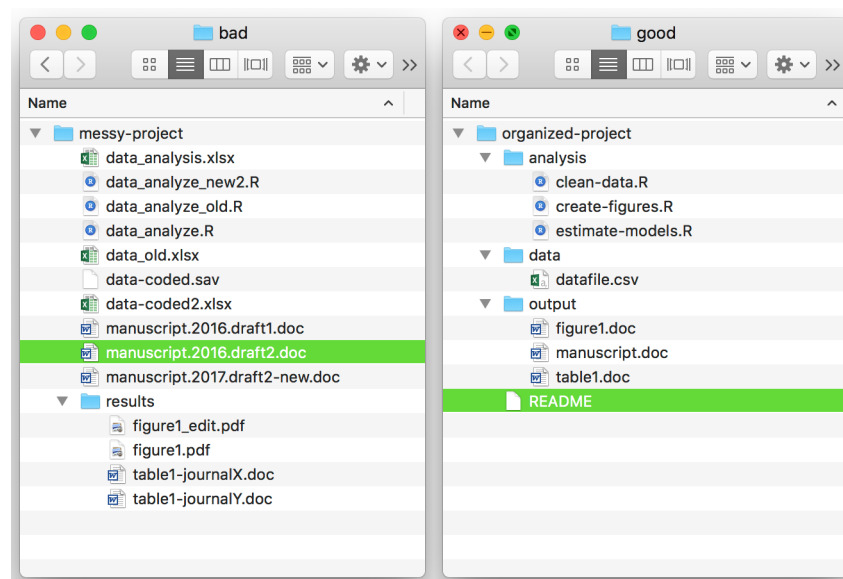


Figure 4. Organizing a project’s files and folders helps potential collaborators (including yourself in the future) to quickly and reliably find the correct files. Some bad practices (left panel) include having multiple versions of each file, increasing the potential for e.g. accidentally using the wrong data for the analysis, or forgetting which version was used in the analysis. Better practices (right panel) organize the data to subfolders with meaningful names, and have one file per purpose. These files are then versioned using Git, eschewing the need for multiple files, thereby reducing potential for mistakes.

Organizing Files and Folders

Implementing reproducibility into the scientific workflow is less time-consuming and effortful if it is planned from the onset of the project, rather than added to the project after all the work has been completed. It is therefore important to organize the project keeping a few key goals in mind (here, we follow guidelines such as the Project TIER recommendations¹¹): The files and folders should have easy to understand names (avoid idiosyncratic naming schemes), and the names should indicate the purposes of the files and folders. We illustrate some good and not-so-good practices in Figure 4.

The first step is to create a home folder for the project. This folder should have an immediately recognizable name, and should be placed somewhere on your computer where you can find it. We call the example project (and therefore its home folder) **git-example**. All the materials related to this project will be placed in sub-folders of the home directory, but for now **git-example** is just an empty folder waiting to be filled with materials. Because the folder structure on a computer is easy to think of as a tree, a project’s home folder—or any folder that has sub-folders—is also known as the *root* directory. In what follows we use the terms home directory and root (directory) interchangeably.

To create a new folder for the Git repository, you first choose an appropriate folder on your computer (such as **User/Documents/**) where you’d like to create the project. You can either use the system’s file navigator (Finder / File Explorer) to create this folder, or use

¹¹<http://www.projecttier.org/tier-protocol/specifications/>

```

Matti [~]$ cd Documents
Matti [~/Documents]$ mkdir git-example
Matti [~/Documents]$ cd git-example
Matti [~/Documents/git-example]$ pwd
/Users/Matti/Documents/git-example
Matti [~/Documents/git-example]$ git init
Initialized empty Git repository in /Users/Matti/Documents/git-example/.git/
Matti [~/Documents/git-example]$

```

Figure 5. Creating and navigating to a folder, and initializing it as a Git repository.

the command line: Navigate to the desired folder by using `cd Documents` to move into the `Documents` folder (assuming it exists in the folder where you currently are). Use `cd ..` to move out of a folder (to its containing folder), if needed.

Initializing a Git Repository

Once you are in the folder where you want to create the project, type `mkdir git-example` to make the `git-example` directory, then `cd git-example` to move into it. You can, of course, also use the point-and-click interface (Finder or File Explorer) to create folders, instead of the `mkdir` command. Once you are in the project's home folder (you can verify where you are by typing `pwd`), you can turn the folder into a Git repository by initializing Git with the `git init` command. These commands are shown in 5

Instead of screenshots, for the rest of the tutorial we present the commands as follows:

```
$ git init
```

In these code listings, each command is preceded by a `$` symbol to indicate that they are inputs to the command line. The command line typically displays this symbol (some versions might use another symbol, such as `>`) on the current line where commands can be entered. Do not type the `$` symbol as part of the command. To help remember the commands, we recommend typing them out, instead of copy-pasting. The `git init` command initializes the folder as a Git repository, and the only change so far has been the addition of a hidden `.git` folder inside `git-example` (and possibly a `.gitattributes` file. Users can ignore these hidden files and folders, however they can be shown with the `ls -la` command.) Now that the folder is initialized as a Git repository, Git monitors any changes within it, and allows you to add and commit these changes (Figure 3.)

Adding a File to Git

Every project (repository) should contain a brief note explaining what the project is about and who to contact. This note is usually called a `readme`, and therefore our first contribution to this project will be a `README` file (this file is so important that it has become standard practice to write it in capital letters). The `README` file should be a plain text `.txt` file (i.e. not created with Microsoft Word) that can be read with a simple text editor¹². You can create this file with any text editor (e.g. the `TextEdit` and `Notepad` apps which allow for creating plain text files are installed by default on Mac and Windows

¹²Plain text has many advantages over proprietary file formats such as Microsoft Word's `.docx` files. Briefly, plain text is both human and computer readable (`.docx` are unreadable without a copy of Microsoft Word),

OS respectively). We created the file with some text in it, and can now see the file in the `git-example` folder, either by using the system's graphical file viewer or with the `ls` command line command.

Because we added this file to a (folder that is a) Git repository, Git is also aware of it. To see what files have changed since the last status change in the repository (there clearly has been only this one), you can ask for Git's **status** (in subsequent code listings, command line input is prepended with `$`, and output is printed without preceding characters):

```
$ git status
On branch master
Initial commit
Untracked files:
  (use "git add <file>..." to include in what will be committed)
  README
nothing added to commit but untracked files present (use "git add" to track)
```

The relevant output returned from executing this command is the "Untracked files:" part. There, Git tells the user that there is an untracked file (`README`) in the repository. To start tracking changes in this file, we **add** it to Git's staging area (Figure 3) by using the command `git add` followed by the file name (i.e. `README`, in our example, the file doesn't have an extension), or `.` which is a shortcut for adding all files with changes to the staging area.:

```
$ git add README
```

We have now added this file to the staging area, and if we are happy with changes to the file's status, we can **commit** the file to Git's history. Commits are essential Git operations: They signify meaningful changes to the repository, and can be later browsed and compared to one another. As such, it is helpful to attach a small message to each commit describing why that commit was made. Here, we've created a `README` file, and our commit command would look as follows:

```
$ git commit -m "Add README file."
```

The quoted text after the `-m` argument is the *commit message*. Entering this command to the command line returns a brief description of the commit, such as how many files changed, and how many characters inside those files were inserted and deleted. The distinction between adding and committing in Git is important: The adding stage allows the user detailed control of what to add to the staging area in preparation of a commit to Git's history. The commit, then, commits the files from the staging area to history. These two operations can be run simultaneously by including the `-a` option to `git commit` (e.g. `git commit -a -m`

is both forward and backward compatible (there will always, and has always been, software capable of reading it), and plain text takes very little space. The file extension (the `.docx` part of a Word file's name, for example) of plain text doesn't matter much, but we recommend using either `.txt` because it is widely recognized, `.md` for markdown syntax, or in the case of a readme file, simply not using any file extension.

265 "Commit message"), but it may be more difficult to control what gets committed with this
266 command, and we therefore recommend beginning users to do `git add` and `git commit` in
267 separate commands.

268 Keeping Track of Changes with Git

269 The `git-example` project (or rather, Git) now keeps track of all and any changes to
270 README. To illustrate, you can change the text in the README file with a text editor,
271 save the file, and then ask for `git status` on the command line:

```
$ git status
On branch master
Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git checkout -- <file>..." to discard changes in working directory)
    modified:   README
no changes added to commit (use "git add" and/or "git commit -a")
```

272 Git can tell that the README file has been modified since the last commit. It is often
273 useful to know exactly *how* a file has changed, before committing it. To view differences to
274 a file not yet committed, use `git diff file`. It shows changes within `file`, line by line,
275 highlighting removed lines of text with red and added lines with green. Once you are happy
276 with the changes, you can repeat the add and commit steps from above to permanently
277 record the current state of the project to Git's history (below we use the `.` shortcut for
278 adding all files with changes¹³):

```
$ git add .
$ git commit -m "Populate README with project description."
```

279 What Does Git Know?

280 The real importance of these somewhat abstract steps becomes apparent when we
281 consider the Git **log**. To reveal the commit log of your repository, call

```
$ git log
```

282 The output of this command shows that each commit is identified with a unique
283 hash code (long alphanumeric string), which we can use to call for further information (see
284 below); an author; a date and time; and a short commit message. Executing `git log` on
285 our example project at this stage returns this:

¹³For the `.` shortcut to work, you must currently be in the project's root directory. If you are not in the root directory, you can use `-A` shortcut instead. However, we recommend that you always ensure that you are in the project's root directory when running any Git commands.

```
commit 60cbe5c9b4a78e500314f791080381030577a035
Author: Matti Vuorre <mv2521@columbia.edu>
Date:   Tue Jun 13 17:20:27 2017 -0400
    Populate README with project description.

commit 16c475023ecbc99446164187eeaaab10647ac550
Author: Matti Vuorre <mv2521@columbia.edu>
Date:   Tue Jun 13 17:14:14 2017 -0400
    Add README file.
```

286 To see what exactly changed in the last commit (latest commits are at the top), you
 287 can call `git show` with the commit's hash code (only relevant parts of output shown below):

```
$ git show 60cbe5c9b4a78e500314f791080381030577a035
Author: Matti Vuorre <mv2521@columbia.edu>
Date:   Tue Jun 13 17:20:27 2017 -0400
    Populate README with project description.
diff --git a/README b/README
--- a/README
+++ b/README
@@ -0,0 +1,2 @@
+# Example Git Project
+This example project illustrates the use of Git.
```

288 This output is a detailed log of the changes in that commit. From top, it lists the
 289 commit's author, date, message, and then the commit's “**diff**”. The diff is a detailed
 290 description of the changes introduced in that commit, explained in more detail below. The
 291 diff's first line tells that the following output is a git diff, and the origin of the changes was
 292 the README file (a/README), and the destination was the same file (b/README):

```
diff --git a/README b/README
```

293 These two file names could be different if the content from a file was moved to another
 294 file with a different name, or if a file was moved. The next two lines expand the information
 295 presented in the previous line:

```
--- a/README
+++ b/README
```

296 The first and second line in this output show the file which received deletions (lines were
 297 removed, ---) and the file which received additions (lines were added, +++). Alternatively,
 298 if a file was moved or renamed, these lines would indicate that. Following these lines, the
 299 output shows where the changes were made:

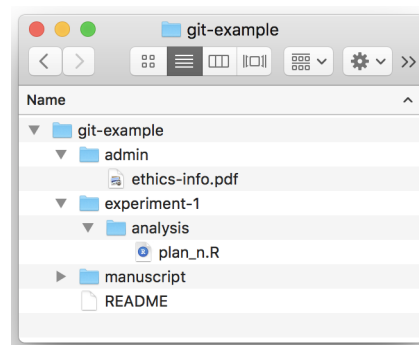


Figure 6. A project with more components.

```
@@ -0,0 +1,2 @@
```

The `-0,0` indicates where in the initial version of the file the changes were made; the first number indicates the first line of changes, the second indicates for how many following lines the changes continued. This output is a little awkward, but because the file was initially empty, the changes must have occurred on the 0th line. The second pair of numbers (`+1,2`) indicates that the added text begins on line 1, and covers two lines of text. The last two lines of the output indicated the changes in the text. The two added lines are prepended with `+` symbols to indicate that these lines were added:

```
+# Example Git Project
+This example project illustrates the use of Git.
```

Although you have now seen the fundamentals of using Git to track the states of (and therefore changes to) a repository, this overly simplistic example doesn't allow full appreciation of the benefits of using Git for version control. To better illustrate Git's functioning, we now fast-forward in the hypothetical example project to a stage where more files and materials have been created.

(Slightly More) Advanced Git

For example, after working for a while on the project, you could have added two files inside well-organized sub-folders of the project. At this point, the project could look like Figure 6.

Running `git status` in this folder reveals that there are two new files (possible empty folders are ignored). That is, since the previous commit, two files have been added to the project; one a .pdf file with some administrative information (`ethics-info.pdf`), the other one is an R (R Core Team, 2017) script file for a pre-planned power analysis (`plan_n.R`). You would probably like to track any changes to the power analysis script, but the ethics information file isn't something that you need to keep track of—at least for this example. It would be laborious to constantly keep ignoring it when committing changes to Git (especially if you'd like to add multiple files simultaneously with `git add .`) Git has an elegant solution to specifying which files to keep track of. Because by default all files are monitored, Git uses a special file for instructing which files are to be *ignored*.

Make Git ignore files. To make Git ignore files, you simply add a plain text file called `.gitignore` to the home folder of the repository. You can use any text editor to create this file¹⁴. Notice that the file is *hidden* (by default, not visible in the OSs file viewer), but can be seen in the command line with the command `ls -la`. Each row of this file should specify a file or a folder (or a regular expression) that Git should ignore. In the current example you could make Git ignore the `admin/` folder entirely (first line in the code listing below), and any file with the `.pdf` extension inside `manuscript/` (second line). The example `.gitignore` file would look like this:

```
admin/
manuscript/*.pdf
```

Re-running `git status` now only shows the `plan_n.R` file and the newly created `.gitignore` file, which is also under version control, naturally. Because there are now two untracked files, which are not specified to be ignored in `.gitignore` (`.gitignore` and `plan_n.R`), and you usually should aim to maintain a clean commit history for the project, you can create two separate commits: One for the `.gitignore` file, and one for the power analysis file.¹⁵

```
$ git add .gitignore
$ git commit -m "Added .gitignore file"
$ git add .
$ git commit -m "Completed power analysis"
```

After this last commit you can, at any time in the future, come back to this commit with `git log` or `git show` and see what was inside the newly created power analysis file when it was first created. For instance, if new information suggests that you should change the assumed effect size in the power analysis, you can simply edit and save the file, then add and commit the changes to Git with a helpful message that logs this important event in Git's history.

This possibility of “rewinding history” is especially useful for files that might go under multiple revisions (manuscripts, analysis files), or if you are interested in when and in what order the files were created—and who created or changed them. One might consider committing a power analysis file to Git as a small personal pre-registration of a part of the research plan. Below, we show some useful ways in which Git can be used to “rewind” the commit history.

Try a new feature. We often find that making some changes to a project didn't have the desired effect: The manuscript ended weaker or the analysis didn't work anymore. Git allows great flexibility in trying new features, then undoing the changes¹⁶. Starting with

¹⁴You can also create this file in the command line by using the `touch` command: `touch .gitignore`. You can then open the file with a text editor to make changes to it.

¹⁵It is entirely up to the user to decide what to commit and when. However, it is best practice to commit often while making incremental changes. Each commit should aim to solve one problem, introduce one new idea, or—more generally—do one thing. This way, when the commit history is reviewed later, it is easy to find and come back to a specific change.

¹⁶A particularly powerful approach for trying new features is **branching**: The project can be duplicated to a new branch and modified, then merged back to the main branch after work on the new feature is

355 an empty staging area, you could start modifying a file (e.g. `plan_n.R`) and later realize
356 that the changes were not good. At this point it is common to press “Undo” in the text
357 editor, but if the file has been saved multiple times or multiple files have been changed, it is
358 difficult to get to the starting point by simply using the “Undo” button. Instead, with Git
359 you can **checkout** the file’s previous version from history. To undo all changes to `plan_n.R`
360 (since the last commit), run

```
$ git checkout experiment-1/analysis/plan_n.R
```

361 Notice that you have to write the full path of the file (relative to the project’s root,
362 see Figure 6) so that Git knows precisely which file you want to checkout from history. With
363 these example operations, we have discussed the main Git operations as outlined in Figure
364 3: Make changes to files, **add** them to the staging area, **commit** to history; **checkout** from
365 history to undo changes. Although these operations are enough to maintain a Git repository,
366 we have only scratched the surface of Git’s functionality.

367 **Undo committed changes.** Another common scenario is one where a user makes
368 changes to a file, adds the changes to the staging area, commits them to Git’s history, and
369 only then realizes that the changes weren’t good. If you have committed changes to a file,
370 and would like to revert back to an older version of the file, you could **checkout** the file’s
371 earlier version, and then commit the older version.¹⁷ For example, suppose you have made
372 bad changes to a file called `file.txt`, and committed the changes to history, and would
373 then like to undo the bad changes by reverting to an older version of the file. View the
374 history with `git log` (you can use the `--oneline` argument for more concise output):

```
$ git log --oneline
4c64f11 Bad changes to file.txt
039d6ff Good changes to file.txt
a73f2ec Add file.txt
```

375 Recall that `git log` returns the most recent changes at the top, and notice that
376 the `--oneline` argument has also made the commit hash codes shorter and thus easier to
377 read and write. Here we can see that commit `039d6ff` has a good version of the file, and
378 subsequent changes in commit `4c64f11` were bad (you would of course not commit “bad
379 changes”, but here the message is informative for clarity). To revert `file.txt` to its good
380 state in commit `039d6ff`, you can use `git checkout hash filename`, which here would
381 be:

```
$ git checkout 039d6ff file.txt
```

382 Now asking for `git status` reveals that `file.txt` has been modified in the working
383 directory (its current state is as it was in commit `039d6ff`). You can now add and

complete—or the new branch can be discarded if the work ended unsatisfactory. Branches are outside the scope of this tutorial, for more information see the Git website (<https://git-scm.com/book/en/v2/Git-Branching-Branches-in-a-Nutshell>).

¹⁷For more information on undoing changes, see <https://www.atlassian.com/git/tutorials/undoing-changes>.

384 commit these changes with `git add file.txt`, then `git commit -m "Undid bad changes`
385 `to file.txt"` (type a commit message suitable for your situation).¹⁸ `git log --oneline`
386 would then show:

```
$ git log --oneline
bcbb123 Undid bad changes to file.txt
4c64f11 Bad changes to file.txt
039d6ff Good changes to file.txt
a73f2ec Add file.txt
```

387 This operation of checking out earlier versions is very useful not only for undoing
388 changes, but for viewing older versions of files as well. However, if you would only like to
389 view past states of the project, instead of reverting / undoing to earlier states of particular
390 files, you can checkout an earlier version of the entire repository, as explained below.

391 **Going to an earlier version of the project.** To return to an earlier state of the
392 project, you can use the `git checkout` command. For example, the display above shows
393 that in commit `a73f2ec`, you had added `file.txt`. If you would like to see the project at
394 that commit, type `git checkout a73f2ec`. This command instantly checks out the file(s)
395 at that point in history, and places them in the working directory where you can view them.
396 This is very helpful if, for example, you would like to quickly run an earlier version of a
397 statistical analysis, which may depend on multiple files. After you have viewed this old
398 version of the project, you can return to the current version with `git checkout master`.

399 Both of these operations—checking out an earlier version of a file or of the entire
400 project—are “safe” in the sense that your project’s history won’t be affected. However,
401 checking out an earlier version of a specific file changes the current state of the project (the
402 current version of the file is temporarily overwritten with the old version), so it is good
403 practice to carefully keep track of the current version of your file before making further
404 commits.

405 Collaborating

406 The true advantages of using Git become apparent when we consider projects with
407 more than one contributor. For example, consider a project where data is collected at
408 multiple sites, and the data files are then saved onto a central server, or shared through a
409 service that automatically merges files from multiple sources (such as the popular Dropbox
410 service). If two or more sites accidentally save a data file with the same name to the server
411 (or Dropbox), and these changes are then automatically merged, the later file will simply
412 overwrite the earlier file. Disaster! Alternatively, consider a data analysis where two or more
413 people work simultaneously on some complicated analysis script, and share their work using
414 a similar system as in the above example. If user A and B are making changes to the same
415 file and user B saves the file, user A’s version of the file will be overwritten. Disaster!

416 Git and other VCSs, on the other hand, were specifically designed to allow (and
417 facilitate) multi-site collaboration on arbitrarily complex projects: Microsoft Windows is

¹⁸If, for some reason, you preferred the latest version after all, you can undo the revert process by `git checkout HEAD file.txt`, instead of adding and committing the older version. This function checks out the current state (“HEAD”) which contained the bad changes.




Figure 7. A diagram illustrating the typical collaborative Git workflow with a remote repository (e.g. GitHub). Verbs indicated in bold text are Git operations and explained in detail in the main text.

developed on a Git platform collaboratively by about 4,000 engineers (Harry, 2017). We believe that Git can be especially helpful in scientific collaboration, a topic to which we turn next.

There are many ways in which a team could collaborate on a Git project; here we focus on a common one, called “centralized workflow”, where one virtual copy of the project is considered central, and all contributors interact with this central copy, instead of directly with each other.

Overview of the Centralized Git Workflow

In this workflow, a central project hosted on a research team’s server, or an online service like GitHub, is used to send and receive information from and to local projects. Because the word “project” can now refer to virtual copies of projects, they are instead commonly referred to as **repositories**. First, one user creates this central repository, then other users **clone** local copies of it (see Box 2). Contributors, including the one who created the central repository, then work on their local projects as detailed above, and after committing, **push** their changes to the central repository. To get changes which other users have pushed to the central repository, contributors **pull** changes from it.

Setting up a centralized Git repository on a research team’s private server is relatively straightforward, but because the details vary from team to team, here we illustrate the centralized workflow using GitHub.

437 GitHub

438 GitHub is one of the 100 most popular website worldwide, and hosts over 60 million
439 software projects with a total of over 20 million users.¹⁹ We chose GitHub for collaboration
440 because it is already the de facto standard among scientists who use VCS, it offers free
441 private repositories (see below), and repositories from GitHub can easily be connected to
442 projects hosted on the Open Science Framework (<https://osf.io>).

443 Anyone wanting to use GitHub must first create a free user account at <https://github.com>.
444 After the user account has been set up, and the user has logged in, creating a new
445 repository at GitHub is self-evident (there is a big green button labeled “New repository” in
446 the user’s home page). You must create the GitHub account with the same email address
447 that you used above when configuring your local Git (or re-configure your local Git to use
448 your the email address that you used to register for GitHub).

449 **Create a new GitHub repository.** Although we describe how to create a new
450 repository, the goal here is to create an online “remote” repository for the running example
451 within this tutorial, so that you can “push” your current local repository to the GitHub
452 remote. Because of this, when you are creating the new repository on GitHub, do not
453 check the “Initialize this repository with a README” (you already created one in the
454 local repository if you followed the steps above on your own computer), and do not add a
455 .gitignore or a license. When creating the repository, you are also asked to enter a name
456 for the repository, which can be anything. For consistency we call the GitHub repository
457 `git-example`—the same name as the local project folder. After entering the name, click
458 “Create repository”. The next step is to link the new remote on GitHub to the local
459 `git-example` repository. To do this, use the following commands (the commands are also
460 visible on GitHub, on a page that opens after you have created the repository):

```
$ git remote add origin https://github.com/username/reponame.git
```

461 Above, `username` and `reponame` are the user’s GitHub user name, and the GitHub
462 repository name. The correct address is visible in the GitHub page that opens after creating
463 the repository. Once you have added the GitHub remote to the local repository, you can
464 verify that the correct address was given with the command `git remote -v`. Once the
465 connection is set up, we can **push** local changes to the GitHub remote:

```
$ git push -u origin master
```

466 The `-u origin master` arguments are only required for the first push, as they set
467 up the connection. Running this command will send your local repository to the GitHub
468 repository. For pushing changes following this initial push, simply type `git push` after
469 adding and committing locally. You have now created the remote central repository, and
470 other users can start contributing to it.

471 **Create a GitHub repository from an existing project.** Alternatively, we may
472 already have a project with several files in it that is not yet a Git repository and that we
473 wish to share with others. To do this, first log in to your GitHub account and create a new

¹⁹<https://github.com/about>.

repository. This can be done by clicking on the large green button called **New** under the **Repositories** tab. On the next page enter the repository name in the text box and choose whether to make the repository public or private. The name of the repository should match the name of the top-level folder of the project you wish to share. We recommend not checking the box for adding a README or adding a `.gitignore` file or license at this point. These can be added after completing the setup. Next open the command line (e.g. the Terminal for Mac users, Git Bash for Windows users) and change the current working directory to the home directory of the project you wish to share.

First initialize the local directory as a Git repository:

```
$ git init
```

Now add all files to the staging area (if there are files or folders that you would not like to include, see the above section on how to make Git ignore files with the `.gitignore` file):

```
$ git add .
```

Commit these files including a comment such as “First commit”:

```
$ git commit -m "First commit"
```

To push the local files to GitHub we follow the same steps as in the previous section where `username` and `reponame` are the user’s GitHub username, and the GitHub repository name.

```
$ git remote add origin https://github.com/username/reponame.git
```

You can then verify the new repository:

```
$ git remote -v
origin https://github.com/username/reponame.git (fetch)
origin https://github.com/username/reponame.git (push)
```

The two lines of output indicate the remote repository from where changes are to be received, and from they are to be sent, respectively. These should usually be the same address. You can then push all files to GitHub with the first commit:

```
$ git push -u origin master
```

Like when sharing a git repository created from scratch, after these steps all users can contribute modifications to the central GitHub repository by pushing changes with `git push` after adding and committing locally.

Contributing to a central (GitHub) repository. Once another team member has set up Git on their own computer, and signed up for GitHub, they need to **clone** the remote GitHub repository onto their local computer. To clone a repository, the new user must first navigate to an appropriate location on the computer where they would like to save the project on their computer, for example their **User/Documents** folder. They do not need to create the project's folder locally before cloning the remote repository. Once the appropriate location is found, cloning will create a new folder for the repository inside this folder. Use `git clone` to clone a repository from a URL:

```
$ git clone https://github.com/username/reponame.git
```

To find out the correct URL to enter to `git clone`, you can navigate your web browser to the repository's GitHub address (e.g. <https://github.com/mvuorre/reproguide-curate> for this tutorial's repository) and click the big green "Clone or download" button; the complete address is in the text box. New contributors can then work on their local copies as detailed in earlier parts of this tutorial; making changes, adding, committing, and pushing. After committing their changes, they can update the status of the central Git repo by pushing their changes to it:

```
$ git push
```

You can also clone GitHub repositories if you are not aiming to contribute to a project. Hosting projects on GitHub is attractive because other people can easily clone (download) your work to build on, and contribute to, your work.

Obtaining other's changes from the central repo. Just as you must manually push your own local changes to the remote repository, you must also obtain others' changes by **pulling** them from the central repo. Pulling is indicated as the first step in the collaborative workflow in Figure 7, because it is important that you start working on the most up to date version of the project (e.g. you don't want to reinvent the wheel or make unnecessary conflicting changes). Before starting to work on your proposed changes, pull the remote changes with:

```
$ git pull
```

The way in which users and their local repositories interact with the central repository by pushing and pulling is the cornerstone of collaboration on GitHub, and thoughtful use of these commands allows for complex workflows without any important code (data, ideas in manuscript, analysis code) ever being overwritten. For instance, prior to restarting work on a project it is often worth checking if any files have changed in the central repository by checking the status of your local repository with `git status` as described above. If someone else has made changes to the central repo, pull them to your local repo before starting your work. It is also worthwhile in committing and pushing changes made after each period of work. However, there is no automatic way for a computer to tell what changes to prioritize: If two or more users have worked on the same code and then attempt to push their changes, it is possible that they have made changes which conflict with each other. If this happens, there is no need to worry; you simply need to know how to resolve the conflict.

Resolving conflicts in collaborative work

Many different types of conflicts may appear in collaborative work, such as multiple users creating files with the same name but different content, or multiple users working on the same code (recall that we use the word code to mean any text written on a computer, e.g. text in a manuscript) and creating changes that conflict with each other. We use this latter situation as an example to explain how to resolve conflicts in collaborative work.

Consider the following scenario. Two collaborators, User A and User B, are working on the same project (they collaborate on a repository on GitHub, with local repositories as detailed above.) At some point, they might be working on the same file—such as writing a manuscript together—and find that they have made changes that conflict with each other. More formally, let's assume that both users are making changes to a file, and User B happens to add and commit his changes locally and push them to the central repository before User A does. When User A then attempts to push her changes to the local repository—and the changes are incompatible with User B's changes—a **conflict** will happen. That is just a natural consequence of two individuals working simultaneously on the same idea, and then writing different code in the same location in the file. When this happens, user A needs to first integrate the latest version of the project from the central repository to her local project, such that it reflects both collaborators' edits, and then push the new “merged” version to the central repository. Let's look at what this workflow entails in a little more detail.

First, let's assume that a collaborator (User B) has made changes to the `README` file in the `git-example` project and pushed the changes to the central repository. For brevity, we only show the first few lines of this file (User B has added the second line):

```
# Example Git Project
Hello world!
This example project illustrates the use of Git.
```

At the same time, User A might have changed her local version of the `README` file to look like this (User A also added to the second line):

```
# Example Git Project
Here are some changes.
This example project illustrates the use of Git.
```

If User A now commits the changes locally, and attempts to push the changes to the central repository, an error will appear

```
$ git add .
$ git commit -m "Some meaningful changes"
$ git push
error: failed to push some refs
hint: Updates were rejected because the remote contains work that you
do not have locally. This is usually caused by another repository
pushing to the same ref. You may want to first integrate the remote
changes (e.g., 'git pull ...') before pushing again.
```

As is usually the case, Git also returns a hint on what to do, which you can follow to successfully resolve the conflict. Git suggests that User A first integrate the remote changes: She needs to first obtain the latest version of the file(s) from the central repository, add the proposed changes to the latest version of the file (the one containing User B's changes), and push that version. She can first obtain the latest changes from the central repository:

```
$ git pull --rebase
```

The `--rebase` argument turns Git into rebasing mode, meaning that you are now applying your local commits on top of the “base” obtained from the remote repository. You can `git pull` without the `--rebase` argument, but that would create an unnecessary commit message and lead to a slightly different workflow²⁰; we recommend using the `--rebase` argument. At this point, Git is in “rebasing” mode, allowing User A to resolve the conflict before pushing her changes. Running `git status` returns (only relevant output shown):

```
$ git status
rebase in progress; onto bada506
You are currently rebasing branch 'master' on 'bada506'.
  (fix conflicts and then run "git rebase --continue")
Unmerged paths:
  (use "git reset HEAD <file>..." to unstage)
  (use "git add <file>..." to mark resolution)
    both modified:   README
no changes added to commit (use "git add" and/or "git commit -a")
```

Git again shows a helpful (if rather jargony) message; the relevant point is that the repository is in rebasing mode (technically, User A is re-basing her current local version (`master`) onto User B's latest contribution, identified with `bada506`). The next step is to manually edit the conflicted file to reflect both User A's and User B's changes. Of course, the user who is doing the rebasing (i.e. merging or resolving the conflict) may simply decide to edit the file to reflect only her changes, completely rejecting User B's changes. When viewed with a text editor, the conflicting `README` file, on User A's computer, now looks like this:

```
# Example Git Project
<<<<<<< HEAD
Hello world!
=====
Here are some proposed changes.
>>>>>>> Some meaningful changes
```

The first line of the file was identical across the two Users' versions of the file, and therefore remains the same. However, after the first line, there is first a line (`<<<<<<< HEAD`) indicating that what follows are the to-be-integrated lines of text. Anything after this line

²⁰For more information, see <https://www.atlassian.com/git/tutorials/comparing-workflows>.

up to the ===== line is the to-be-integrated text from the central repository. We can see that this is simply the line of text that User B created (“Hello world!”). After the separating line (=====) are User A’s local changes, followed by those changes’ commit message (“Some meaningful changes”) prepended with a >>>>>>. User A can then edit this file however she chooses, using the tags to help her see which are her lines of code (text), and which are User B’s. For example, User A may integrate the changes to make the file look like this:

```
# Example Git Project
Here are some proposed changes: Hello world!
This example project illustrates the use of Git.
```

Then, User A can save the file and add the changes with `git add README`. Importantly, because Git is in rebasing mode (which can be aborted with `git rebase --abort` to reject the central repository’s changes and return User A back to her latest local version), User A should not commit but instead needs to complete the rebasing with `git rebase --continue`.

```
$ git rebase --continue
Applying: Some meaningful changes
```

This command returns a reminder telling User A which local commit is being applied on top of the changes she pulled from the central repository, and then returns Git to its normal mode from rebase mode. The final step is then to use `git push` to send the local changes to the central repository.

How these potential conflicts appear depends on how users collaborate with one another, and a detailed explanation of all potential scenarios is outside the scope of this tutorial²¹. Most importantly, even in the event of conflicts, all committed changes are saved in Git’s history and can be retrieved, so experimenting with different approaches to resolving conflicts is safe.

Private or public collaboration?

By default, all GitHub repositories are public: Anyone with an internet connection can use their web browser to inspect the contents of your repository, or even clone it to their computer. This may sound unfamiliar to researchers used to working more privately, and clearly necessitates planning and thought with respect to issues such as data privacy and sharing sensitive materials. However, for many projects—including the writing of this tutorial—we see very few downsides to working “in the open”.

There are two alternatives to working in a public GitHub repository: One, which we won’t cover here, is to not use GitHub but instead place the central repository on the research team’s private server. The second option is to make the repository private on

²¹Covering all different types of file conflicts is outside the scope of this tutorial. Although the instructions provided herein will help in most common use case scenarios, readers can refer to the following websites for more information: <https://help.github.com/articles/resolving-a-merge-conflict-using-the-command-line/> and <https://www.atlassian.com/git/tutorials/comparing-workflows>. You can also resolve conflicts on GitHub (<https://help.github.com/articles/resolving-a-merge-conflict-on-github/>), and the GitHub customer service is very responsive to users’ help requests, which can include questions on code conflicts.

611 GitHub (this can be done when the repository is first created or afterwards by clicking
612 Settings on the repository’s website). Private repositories, and their contents, are only
613 accessible to invited team members, and are therefore ideal for small teams who would like
614 to work without revealing their master plans to the public just yet. For example, you might
615 initially choose to work in a private repository, and only make it public once you feel the
616 material is mature enough for public consumption.

617 To make a GitHub repository private, navigate to the repository’s website with a web
618 browser, and click “Settings”, then “Make this repository private”. Once one user has set
619 the central GitHub repository to private mode, anyone wishing to clone, push, pull, or view
620 the repo must provide their GitHub username and password. Only if they match an invited
621 team members username and password can the user access the repository. At the time of
622 this writing, GitHub offers five private repositories for free²².

623 Using Git with a Graphical Interface

624 Above, we have detailed how to use Git from the computer’s command line, but not all
625 users are comfortable—although there is no reason not to be!—with this mode of interacting
626 with their computers. We chose to introduce Git from the command line because eventually
627 users might want to use it for more complicated operations (and it is required for setting the
628 user’s information). However, there are various graphical user interfaces (GUIs) available
629 for Git which allow point-and-click interacting with Git.

630 R Studio

631 In this section, we introduce how to use the popular R Studio²³ IDE (RStudio Team,
632 2016) for the R programming language (R Core Team, 2017) for interacting with Git. An
633 IDE is an interface that bundles together many tools for software development—some users
634 may be familiar with R Studio for conducting statistical analyses. Version control is an
635 essential feature of software development, and R Studio provides a good GUI within its IDE
636 for controlling Git²⁴.

637 For this example, we will begin a completely new project using only R Studio, and
638 assume that readers are familiar with the concepts from the above tutorial. Once downloaded
639 and installed, it can be used for accessing the R programming language, and for many
640 operations involved in curating research assets, including a GUI for Git.

641 Managing projects (and Git repositories) with R Studio is centered on the idea of R
642 Projects. To start a new project with R Studio, open the R Studio application, and click
643 File -> New Project. This brings up a dialog (left panel in Figure 8) asking whether to
644 create a project in a new directory, existing directory, or checkout an existing project from a
645 version control repository. Here, we create the project a new directory. We then choose “New
646 Project” in the following screen (middle panel in Figure 8). Then we’ll give a name to the

²²To obtain the free repositories, fill out the request form at https://education.github.com/discount_requests/new. Students with .edu email addresses, can obtain unlimited free repositories at <https://education.github.com/pack>.

²³R Studio is a free and open source IDE, works on Windows, Mac, and Linux operating systems, and can be downloaded from the project’s website at <https://www.rstudio.com/products/rstudio/download/>.

²⁴Although in this example we refer to the R programming language, Git can be used through the R Studio GUI even if you do not use R (or any other programming language.)




Figure 8. Creating a project in R Studio.




Figure 9. R Studio's Git tab for a newly created project. The Git tab in R Studio has buttons for Git commands covered in this tutorial, such as viewing Git's history, diffs, and commits.

project's home folder and choose where to save it on the computer. Importantly, we'll also check the "Create a git repository" box (right panel in Figure 8), which will automatically set up a new repository for the project. Clicking "Create project" creates the folder in the specified location, and two files inside the project's main folder.

One of these files is `.gitignore` which we covered above. The other file is an `.Rproj` file, which indicates that the folder is the home folder for an R (Studio) project. Users don't interact with this file directly, but if opened, it is a plain text file containing the project's settings (these can be modified through Tools -> Project Options in R Studio).

Using Git with R Studio

Once the R Project has been created, R Studio has a "Git" tab in the top-right panel of the GUI. At first, this tab shows the two new files in the repository, and some buttons with familiar looking names ("Commit", "Push", etc.)

Add and commit changes. To mark this milestone of creating a project, you can commit all the changes so far (new files, new project) to Git. The workflow is exactly the same as when using Git from the terminal, but instead of text commands, you can now use clickable buttons. To begin committing these changes, click on "Commit" in the Git tab. This brings up another window where you can select files to add into the staging area (the files in Figure 10 are staged by checking the boxes on the names' left side). The green "A" indicates that these files are new. Once all the desired files (here we chose both) are staged, you can write a short commit message in the appropriate text box, and click the "Commit" button. The text highlighted in green in the bottom part of Figure 10 indicates the lines of text that were changed in the selected file (`.gitignore`). These lines are all green because the file is new to Git and therefore each change is an addition (the files are unimportant R project files which users can safely ignore.)

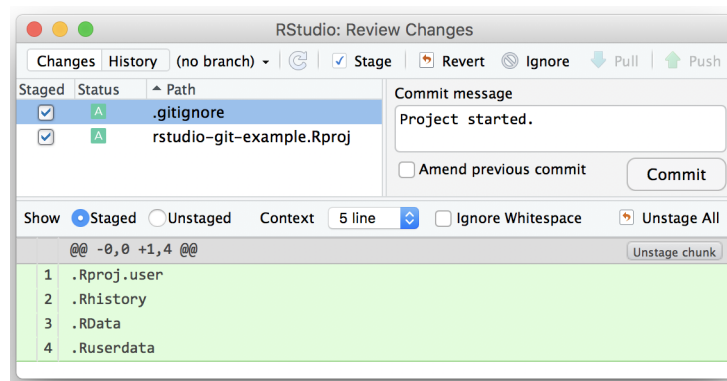


Figure 10. Creating a Git commit with R Studio.

Add a remote GitHub repository. When creating the R Project, there was an option for creating the project from an existing Git repository. Because we didn't do this for the current example, we must now manually instruct Git that this project should have a remote GitHub repository. This is done exactly as above, using the command line functions. After you've created a new repository on GitHub, navigate to the project's folder with the command line shell (or you can click on "More" in R Studio's Git tab, and then "Shell", which opens a new shell window in the correct location), and execute the following commands (but replace `username` and `reponame` with the correct user and repository names):

```
$ git remote add origin https://github.com/username/reponame.git
$ git push -u origin master
```

After executing these commands, you can use all Git and GitHub features from R Studio. This is especially useful because the R Studio IDE offers a complete environment for project management, data analysis, and manuscript preparation with the R Markdown and knitr R packages (Allaire et al., 2016; Xie et al., 2016). Psychologists will be especially interested in the papaja package for creating APA formatted manuscripts (Aust & Barth, 2016). The source code of this manuscript, which was prepared with the papaja package, can be found at <https://github.com/mvuorre/reproguide-curate>.

Git Summary

Box 2 describes the main Git commands and their purported use in roughly the order in which they would be used in a typical workflow. Help on how to use each command is available by appending `--help` to each command (e.g. `git commit --help`). In Box 3 we give links to online materials that we have found particularly helpful in learning Git and teaching it to others.

Box 2. Main Git commands

git init: Initialize a local repository.

- Create a local Git repository inside a folder on your computer. In effect, this command turns the local folder into a Git repository.

git clone URL: Obtain a copy of a remote (GitHub) repository to local computer.

- Create a local Git repository from a pre-existing remote Git repository, such as one hosted on GitHub.

git add file: Add changes in `file` to Git's staging area.

- Once a file has been edited or created in the local project, it needs to be added to Git's staging area. Multiple files can be added simultaneously using the `.` shortcut (**git add .**). Once the files have been added, they can be saved to Git's history.

git commit: Save changes from staging area to Git's history.

- A commit creates a snapshot of the project's current state by adding the most recent changes (which have been added to the staging area) to Git's history. Commits should be accompanied with short, descriptive messages to describe the changes introduced in that commit (**git commit -m "Commit message"**)

git log: Show Git's history.

- Git's log is a list of all the commits made in the current repository, and as such is a description of the project's history. The history can be very long and span multiple pages; to move forward one page, press the space bar, or "q" to quit viewing the log. To view the log in a more concise form, add the `-oneline` argument (**git log -oneline**).

git push: Send locally committed changes to remote repository.

- In the centralized workflow for collaboration as described in this tutorial, collaborators work on their own local repositories, and periodically send the changes that they have made to the central (remote) repository.

git pull: Receive changes from the remote repository.

- To keep the local repository up-to-date with other collaborators' changes, users must periodically pull changes from the remote (central) repository. It is good practice to pull the most up-to-date version of the project every time before making changes, to avoid unnecessary conflicts.

692

693 **Deleting a Git Repository**

694 Finally, users may sometimes choose to delete their Git repositories. Deleting a project
695 is as simple as moving the containing folder(s) to Trash (Recycle Bin on Windows), which
696 also deletes the Git project (the Git project is contained in a hidden `.git` file in the project's
697 home folder.) If you wish to only delete the Git repository, but keep the project itself, you
698 can delete the `.git` folder. Because the folder is hidden, it will not show up in the default
699 graphical file explorer. You can remove this folder using the command line by navigating to
700 the project's root folder, and using the following command:

```
rm -rf .git
```

701 We recommend caution with this operation, as it will permanently delete the `.git`
702 folder, which may contain important information. To verify that the folder was deleted, you

can list all the files and folders in the current working directory, including hidden ones, with the following command:

```
ls -la
```

To delete a repository on GitHub, use your internet browser to navigate to the repository’s GitHub address, click Settings, then “Delete this repository”. Be aware that if anyone has cloned this project to their computer, you cannot delete their cloned versions.

Discussion

We have presented an introductory tutorial on using the Git version Control System for curating and collaborating on research assets in behavioral sciences. The essential Git workflow includes adding and committing incremental changes to a version controlled repository, which can then be collaboratively worked on by many researchers through the online GitHub platform. Although this tutorial includes enough material to get started, Git (and GitHub) is a vast ecosystem with great opportunities, some of which are further discussed by Perez-Riverol and colleagues (2016). For example, the concept of “Born open data”, where research data is automatically posted online upon collection, is made very easy with the Git + GitHub workflow (Rouder, 2016). The challenges to reproducibility are many, and they have only recently received the targeted attention they deserve in the collaborative effort to improve the reliability of empirical sciences. Curating research assets and focusing on the practical aspects of the scientific workflow is important for ensuring the continuity of one’s work, and for efforts toward a cumulative and reliable science.

Box 3. Further resources for learning Git

- [Basic VCS workflow](https://www.git-tower.com/blog/workflow-of-version-control), an infographic explaining how VCS works (<https://www.git-tower.com/blog/workflow-of-version-control>).
- [GitHub’s Git cheat sheet](https://services.github.com/resources/cheatsheets/) is available in multiple languages and contains the most used Git commands (<https://services.github.com/resources/cheatsheets/>).
- [GitHub’s Git glossary](https://help.github.com/articles/github-glossary/) describes Git’s most common commands (<https://help.github.com/articles/github-glossary/>).
- [TryGit](https://try.github.io), an interactive website for learning the basics of Git (<https://try.github.io>).
- [Git + GitHub](http://r-pkgs.had.co.nz/git.html) in an R programming context (<http://r-pkgs.had.co.nz/git.html>).
- [Pro Git book](https://git-scm.com/book/en/v2), a complete manual of Git (<https://git-scm.com/book/en/v2>).

Author contributions

MV and JPC designed the format of the tutorial, MV drafted the manuscript, MV and JPC wrote and approved the final version of the manuscript for submission.

Acknowledgments

We thank Travis Riddle and Judy Xu for feedback on earlier drafts of this manuscript. This work was supported, in part, by Institute of Education Science grant (R305A150467). The authors are solely responsible for the content of this article.

References

- Allaire, J. J., Cheng, J., Xie, Y., McPherson, J., Chang, W., Allen, J., ... Hyndman, R. (2016). Rmarkdown: Dynamic Documents for R (Version 1.3). Retrieved from <https://cran.r-project.org/web/packages/rmarkdown/index.html>
- Aust, F., & Barth, M. (2016). *Papaja: Create APA manuscripts with RMarkdown*. Retrieved from <https://github.com/crsh/papaja>
- Eglen, S. J., Marwick, B., Halchenko, Y. O., Hanke, M., Sufi, S., Gleeson, P., ... Poline, J.-B. (2017). Toward standard practices for sharing computer code and programs in neuroscience. *Nature Neuroscience*, 20(6), 770–773. doi:10.1038/nn.4550
- Harry, B. (2017). *The largest git repo on the planet*. Retrieved June 20, 2017, from <https://blogs.msdn.microsoft.com/bharry/2017/05/24/the-largest-git-repo-on-the-planet/>
- Ihle, M., Winney, I. S., Krystalli, A., & Croucher, M. (2017). Striving for transparent and credible research: Practical guidelines for behavioral ecologists. *Behavioral Ecology*. doi:10.1093/beheco/arx003
- Markowetz, F. (2015). Five selfish reasons to work reproducibly. *Genome Biology*, 16(1), 274. doi:10.1186/s13059-015-0850-7
- McMillan, R. (2005, April 20). After controversy, Torvalds begins work on "Git". *PC World*. Retrieved from https://www.pcworld.idg.com.au/article/129776/after_controversy_torvalds_begins_work_git_/
- Munafò, M. R., Nosek, B. A., Bishop, D. V. M., Button, K. S., Chambers, C. D., Sert, N. P. du, ... Ioannidis, J. P. A. (2017). A manifesto for reproducible science. *Nature Human Behaviour*, 1, 0021. doi:10.1038/s41562-016-0021
- Perez-Riverol, Y., Gatto, L., Wang, R., Sachsenberg, T., Uszkoreit, J., Leprevost, F. da V., ... Vizcaíno, J. A. (2016). Ten Simple Rules for Taking Advantage of Git and GitHub. *PLOS Computational Biology*, 12(7), e1004947. doi:10.1371/journal.pcbi.1004947
- R Core Team. (2017). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from <https://www.R-project.org/>
- Rouder, J. N. (2016). The what, why, and how of born-open data. *Behavior Research Methods*, 48(3), 1062–1069. doi:10.3758/s13428-015-0630-z
- RStudio Team. (2016). *RStudio: Integrated Development Environment for R*. Boston, MA: RStudio, Inc. Retrieved from <http://www.rstudio.com/>
- Vanpaemel, W., Vermorgen, M., Deriemaeker, L., & Storms, G. (2015). Are We Wasting a Good Crisis? The Availability of Psychological Research Data after the Storm. *Collabra: Psychology*, 1(1). doi:10.1525/collabra.13
- Wicherts, J. M., Borsboom, D., Kats, J., & Molenaar, D. (2006). The poor availability of psychological research data for reanalysis. *American Psychologist*, 61(7), 726–728. doi:10.1037/0003-066X.61.7.726
- Xie, Y., Vogt, A., Andrew, A., Zvoleff, A., Simon, A., Atkins, A., ... Foster, Z. (2016). Knitr: A General-Purpose Package for Dynamic Report Generation in R (Version 1.15.1). Retrieved from <https://cran.r-project.org/web/packages/knitr/index.html>